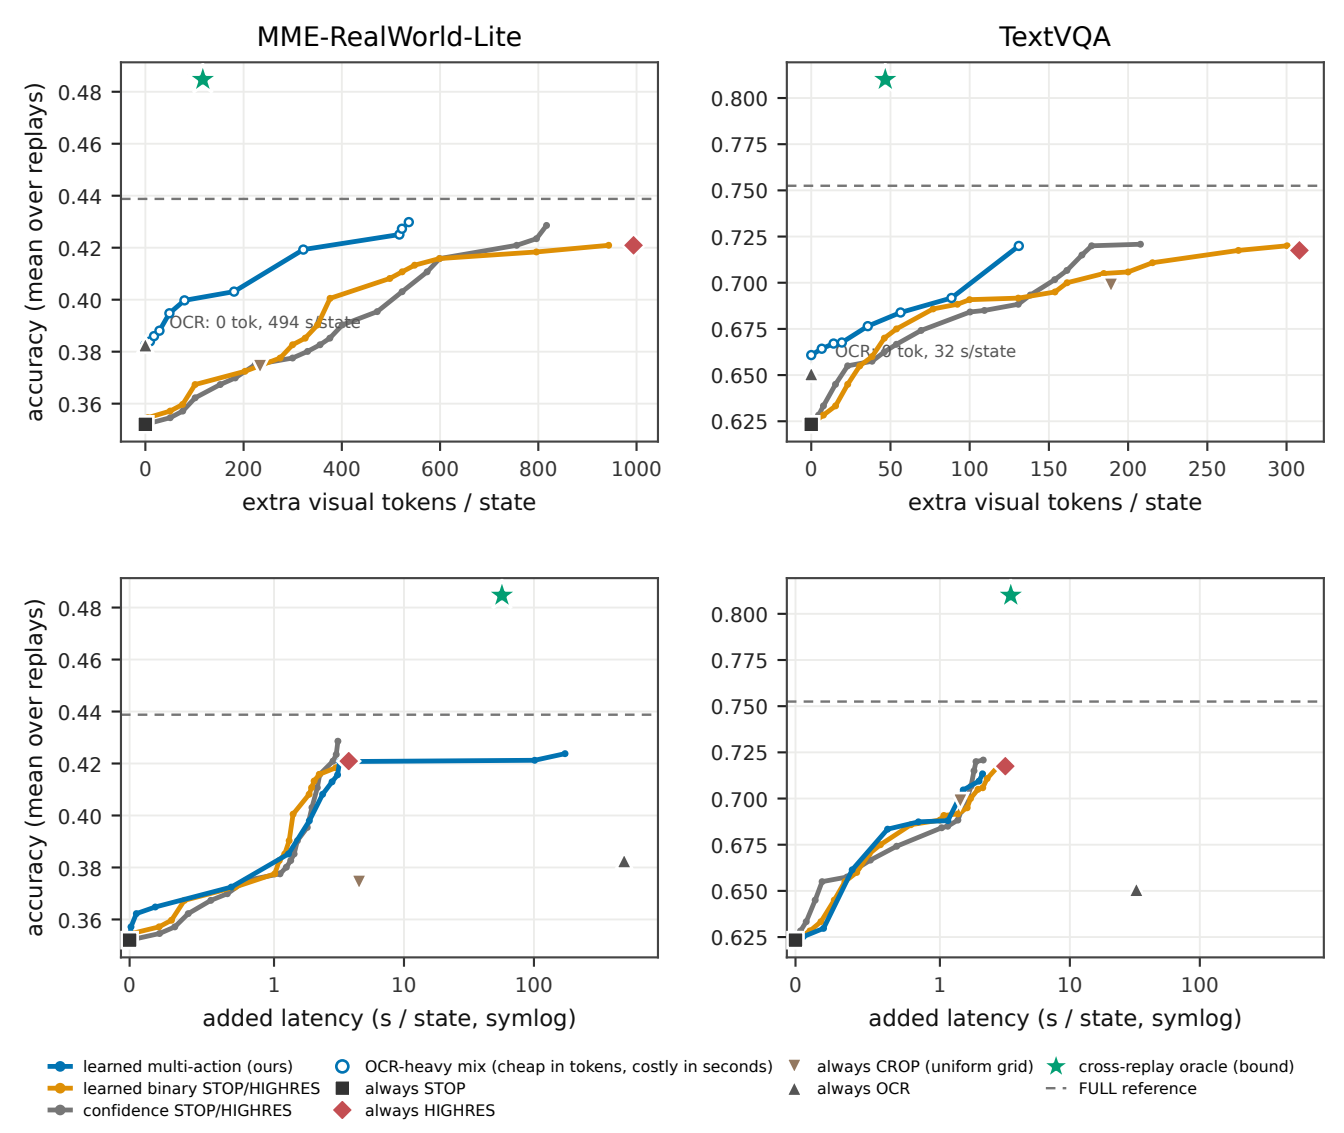
MME-RealWorld-Lite
TextVQA
accuracy (mean over replays)
extra visual tokens / state
added latency (s / state, symlog)
OCR: 0 tok, 494 s/state
OCR: 0 tok, 32 s/state
learned multi-action (ours)
learned binary STOP/HIGHRES
confidence STOP/HIGHRES
OCR-heavy mix (cheap in tokens, costly in seconds)
always STOP
always HIGHRES
always CROP (uniform grid)
always OCR
cross-replay oracle (bound)
FULL reference